

# Preparing for fine-tuning

INTRODUCTION TO LLMS IN PYTHON




Jasmin Ludolf

Senior Data Science Content Developer,  
DataCamp

# Pipelines and auto classes


## Pipelines: `pipeline()`

- Streamlines tasks
- Automatic model and tokenizer selection
- Limited control




## Auto classes (`AutoModel` class)

- Customization
- Manual adjustments
- Supports fine-tuning




# LLM lifecycle



## Pre-training

- Broad data
- Learn general patterns

# LLM lifecycle



## Pre-training

- Broad data
- Learn general patterns

## Fine-tuning

- Domain specific
- Specialized tasks

# Loading a dataset for fine-tuning

```
from datasets import load_dataset  
  
train_data = load_dataset("imdb", split="train")  
train_data = data.shard(num_shards=4, index=0)  
  
test_data = load_dataset("imdb", split="test")  
test_data = data.shard(num_shards=4, index=0)
```

- `load_dataset()` : loads a dataset from Hugging Face hub
  - `imdb`: review classification

# Auto classes

```
from transformers import AutoModel, AutoTokenizer  
from transformers import AutoModelForSequenceClassification  
  
model = AutoModelForSequenceClassification.from_pretrained("bert-base-uncased")  
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
```

# Tokenization

```
from transformers import AutoTokenizer, AutoModelForSequenceClassification
from datasets import load_dataset

train_data = load_dataset("imdb", split="train")
train_data = data.shard(num_shards=4, index=0)
test_data = load_dataset("imdb", split="test")
test_data = data.shard(num_shards=4, index=0)

model = AutoModelForSequenceClassification.from_pretrained("bert-base-uncased")
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

# Tokenize the data
tokenized_training_data = tokenizer(train_data["text"], return_tensors="pt", padding=True, truncation=True,
                                     max_length=64)

tokenized_test_data = tokenizer(test_data["text"], return_tensors="pt", padding=True, truncation=True,
                                max_length=64)
```

# Tokenization output

```
print(tokenized_training_data)
```

```
{'input_ids': tensor([[ 101,  1045, 12524,  1045,  2572,  8025,  1011,  3756,
2013,  2026, 2678,  3573,  2138,  1997,  2035,  1996,  6704,  2008,  5129,  2009,
2043,  2009, 2001,  2034,  2207,  1999,  3476,  1012,  1045,  2036, ...]
```

# Tokenizing row by row

```
def tokenize_function(text_data):  
    return tokenizer(text_data["text"], return_tensors="pt", padding=True, truncation=True, max_length=64)  
  
# Tokenize in batches  
tokenized_in_batches = train_data.map(tokenize_function, batched=True)  
  
# Tokenize row by row  
tokenized_by_row = train_data.map(tokenize_function, batched=False)
```

```
Dataset({  
    features: ['text', 'label', 'input_ids', 'token_type_ids', 'attention_mask'],  
    num_rows: 1563  
})
```

# Subword tokenization

- Common in modern tokenizers
- Words split into meaningful sub-parts

Unbelievably

# Subword tokenization

- Common in modern tokenizers
- Words split into meaningful sub-parts



The word "Unbelievably" is displayed in a large, bold, black sans-serif font. It is composed of four overlapping green rounded rectangles. The first rectangle covers the letters "U" and "n". The second rectangle covers the letters "b", "e", "l", "i", "e", "v", and "a". The third rectangle covers the letters "v", "a", "b", and "l". The fourth rectangle covers the letter "y". The overlapping nature of the rectangles creates a sense of depth and highlights the individual subwords within the word.

Unbelievably

# **Let's practice!**

**INTRODUCTION TO LLMS IN PYTHON**

# Fine-tuning through training

INTRODUCTION TO LLMS IN PYTHON



Jasmin Ludolf

Senior Data Science Content Developer,  
DataCamp

# Training Arguments

```
from transformers import Trainer,  
TrainingArguments  
  
training_args = TrainingArguments(  
    output_dir='./finetuned',  
    evaluation_strategy="epoch",  
    num_train_epochs=3,  
    learning_rate=2e-5,  
)  
 )
```

- `TrainingArguments()` : customize training settings
- See documentation for all parameters
- Values depend on use, dataset, speed
- `output_dir` : output directory
- `eval_strategy` : when to evaluate "epoch", "steps", or "none"
- `num_train_epochs` : number of training epochs
- `learning_rate` : for optimizer

# Training Arguments

```
from transformers import Trainer,  
TrainingArguments  
  
training_args = TrainingArguments(  
    output_dir='./finetuned',  
    evaluation_strategy="epoch",  
    num_train_epochs=3,  
    learning_rate=2e-5,  
    per_device_train_batch_size=8,  
    per_device_eval_batch_size=8,  
    weight_decay=0.01,  
)
```

- `per_device_train_batch_size` and `per_device_eval_batch_size` define the batch size
- `weight_decay` : applied to the optimizer to avoid overfitting

# Trainer class

```
from transformers import Trainer,  
TrainingArguments  
  
training_args = TrainingArguments(...)  
  
trainer = Trainer(  
    model=model,  
    args=training_args,  
    train_dataset=tokenized_training_data,  
    eval_dataset=tokenized_test_data,  
    tokenizer=tokenizer  
)  
  
trainer.train()
```

- `model` : the model to fine-tune
- `args` : the training arguments
- `train_dataset` : the data used for training
- `eval_dataset` : the data used for evaluation
- `tokenizer` : the tokenizer

Number of training loops: Dataset size, `num_train_epochs`, `per_device_train_batch_size` and `per_device_eval_batch_size`

# Trainer output

```
{'eval_loss': 0.398524671792984, 'eval_runtime': 33.3145, 'eval_samples_per_second': 46.916,  
'eval_steps_per_second': 5.883, 'epoch': 1.0}  
{'eval_loss': 0.1745782047510147, 'eval_runtime': 33.5202, 'eval_samples_per_second': 46.629,  
'eval_steps_per_second': 5.847, 'epoch': 2.0}  
{'loss': 0.4272, 'grad_norm': 15.558795928955078, 'learning_rate': 2.993197278911565e-06,  
'epoch': 2.5510204081632653}  
{'eval_loss': 0.12216147780418396, 'eval_runtime': 33.2238, 'eval_samples_per_second': 47.045,  
'eval_steps_per_second': 5.899, 'epoch': 3.0}  
{'train_runtime': 673.0528, 'train_samples_per_second': 6.967, 'train_steps_per_second': 0.874,  
'train_loss': 0.40028538347101533, 'epoch': 3.0}  
TrainOutput(global_step=588, training_loss=0.40028538347101533, metrics={'train_runtime': 673.0528,  
'train_samples_per_second': 6.967, 'train_steps_per_second': 0.874,  
'train_loss': 0.40028538347101533, 'epoch': 3.0})
```

# Using the fine-tuned model

```
new_data = ["This is movie was disappointing!", "This is the best movie ever!"]  
  
new_input = tokenizer(new_data, return_tensors="pt", padding=True, truncation=True, max_length=64)  
  
with torch.no_grad():  
    outputs = model(**new_input)  
  
predicted_labels = torch.argmax(outputs.logits, dim=1).tolist()  
  
label_map = {0: "NEGATIVE", 1: "POSITIVE"}  
for i, predicted_label in enumerate(predicted_labels):  
    sentiment = label_map[predicted_label]  
    print(f"\nInput Text {i + 1}: {new_data[i]}")  
    print(f"Predicted Label: {sentiment}")
```

# Fine-tuning results

Input Text 1: This is movie was disappointing!

Predicted Sentiment: NEGATIVE

Input Text 2: This is the best movie ever!

Predicted Sentiment: POSITIVE

# Saving models and tokenizers

```
model.save_pretrained("my_finetuned_files")
```

```
tokenizer.save_pretrained("my_finetuned_files")
```


```
# Loading a saved model
model = AutoModelForSequenceClassification.from_pretrained("my_finetuned_files")
tokenizer = AutoTokenizer.from_pretrained("my_finetuned_files")
```

# **Let's practice!**

**INTRODUCTION TO LLMS IN PYTHON**

# Fine-tuning approaches


INTRODUCTION TO LLMS IN PYTHON



Jasmin Ludolf


Senior Data Science Content Developer,  
DataCamp

# Fine-tuning




# Full fine-tuning

- The entire model weights are updated
- Computationally expensive



# Partial fine-tuning

- Some layers are fixed
- Only task-specific layers are updated



# Transfer learning

- A pre-trained model is adapted to a different but related task
- Leverages knowledge from one domain to a related one

## Transfer learning

Partial fine-tuning

Full fine-tuning

Zero-shot  
learning

One-shot learning  
Few-shot learning

[...]

# N-shot learning

- Zero-shot learning: no examples
- One-shot learning: one example
- Few-shot learning: several examples

# One-shot learning

```
from transformers import pipeline

generator = pipeline(task="sentiment-analysis", model="distilbert-base-uncased-finetuned-sst-2-english")

input_text = """
Classify the sentiment of this sentence as either Positive or Negative.

Example:
Text: "I'm feeling great today!" Sentiment: Positive
Text: "The weather today is lovely." Sentiment:
"""

result = generator(input_text, max_length=100)
print(result[0]["label"])
```

POSITIVE

# **Let's practice!**

**INTRODUCTION TO LLMS IN PYTHON**